

AI First Principles

AI isn't coming for jobs, it's coming for the bureaucracy that makes work miserable. But bolting AI onto broken processes just scales the mess faster. Eliminating organizational dysfunction demands rethinking both how you design technology and how you rebuild the operations around it.

The AI First Principles guide teams charged with operationalizing AI. They emerged from decades of building for users, understanding what breaks and why, and codifying what holds under pressure.

AI Inherits Messiness

AI learns from people. Therefore AI systems, like people, are inconsistent and operate more effectively with structure. Trying to engineer them to operate like deterministic code will result in system failures. Variation is inevitable, not accidental. Define what's prohibited over what's required.

AI Fails Silently

AI accumulates errors across thousands of interactions before patterns become visible. Traditional systems failed loudly with clear signals; AI fails quietly on repeat. Build feedback loops over post-mortems.

People Own Objectives

AI shouldn't be used in place of human discernment, judgment, or taste. When AI makes mistakes or causes harm, a person should be held accountable, not the algorithm. Name the owner.

Deception Destroys Trust

When AI pretends to be human, people cannot calibrate their expectations, recognize its limitations, and protect themselves from its failures. Make AI obvious, not hidden.

Individuals First

AI industrializes manipulation by personalizing it at scale. Mass persuasion becomes individual manipulation. Build tools that people control, not tools that control people. Prioritize individual agency above efficiency, profit, or convenience.

Build from User Experience

Without input from end users, AI solutions won't solve real problems. People wrestling with system failures are the ones qualified to design system futures. Design systems from lived experience, not distant observation.

Discovery Before Disruption

Changing systems that aren't understood creates unpredictable failures. Redundancies prevent edge cases and manual steps catch exceptions. Existing inefficiencies are containers of knowledge. Identify purpose before simplifying.

Ambiguity Is Wisdom

Concealing ambiguity removes opportunities for critical judgment. Not all decisions are binary yes/no. Wisdom lives in gray areas. AI produces probabilities that demand judgment, not facts that replace it. Surface the probabilities.

Reveal the Invisible

There's a wealth of ignorance hiding in document theater. Expose what you don't yet understand by learning how to articulate it. Pursue what is hard to explain.

Iterate Towards What Works

Grand plans commit to solutions without validating problems. Iteration tests assumptions and measures impact, revealing what works gradually over time. Inherited practices carry outdated logic that meetings can't expose. Learn by doing, not planning.

Decompose Incrementally

Legacy systems carry too much technical debt to replace and are too brittle to automate. AI systems should allow isolated components to decompose naturally. Dismantle legacy complexity piece-by-piece.

Justify Resource Consumption

AI makes it trivially easy to waste resources. What costs pennies to create can cost millions to run. The friction that once prompted resource consideration has vanished. The resources remain real: energy, water, compute, time. Optimize the ratio of value per resource spent.

AI First Principles License

Copyright (c) 2025 AI First Principles (aifirstprinciples.org)

AI First Principles is licensed under the Creative Commons Attribution 4.0 International License (CC BY 4.0). To view a copy of this license, visit <https://creativecommons.org/licenses/by/4.0/>.

Attribution:

When using, sharing, or adapting the AI First Principles, you must provide clear attribution to "AI First Principles" and include a link to <https://aifirstprinciples.org>.

example: For the complete AI First Principles including updates and additional resources, visit aifirstprinciples.org.

Requests For Support

We encourage, but do not require, the following to support the AI First Principles project:

1. **Notification:** Let us know how and where you have used the AI First Principles by emailing info@aifirstprinciples.org.
2. **Contribution:** If you have suggestions for changes or improvements, please share them via our GitHub repository at <https://github.com/aifirstprinciples/AI-First-Principles>.
3. **Financial Support:** Consider supporting the promotion and development of AI First Principles through our Patreon at <https://www.patreon.com/c/AIFirstPrinciples>.